

# **EGR 5110: Homework #6**

Due on May 15th, 2024 at 11:59pm

*Professor Nissenson*

**Francisco Sanudo**

## Contents

<b>Background: Optimization and Gradient Ascent with Backtracking Line Search</b>	<b>3</b>
Objective Function . . . . .	3
Gradient Ascent . . . . .	3
Backtracking Line Search . . . . .	3
Armijo Condition . . . . .	3
Armijo Condition vs. Exact Line Search . . . . .	3
<b>Gradient Ascent Algorithm in MATLAB</b>	<b>4</b>
Inputs . . . . .	4
Outputs . . . . .	4
Optimization Process . . . . .	4
Utility Functions . . . . .	4
Recommendations . . . . .	5
<b>Simulations</b>	<b>5</b>
Objective function and Optimization Setup . . . . .	5
Scenarios . . . . .	5
Intepretation of Results . . . . .	6
<b>MATLAB Code</b>	<b>9</b>

## Background: Optimization and Gradient Ascent with Backtracking Line Search

Optimization is a fundamental problem in mathematics and computer science, where the goal is to find the best solution (maximum or minimum) of an objective function within a given domain. In numerical methods, optimization algorithms are employed to efficiently navigate through the solution space and locate optimal points.

### Objective Function

Consider an objective function  $f(x, y)$  representing a scalar-valued function of two variables  $x$  and  $y$ . The goal of optimization is to maximize  $f(x, y)$  with respect to  $(x, y)$  within a specified region.

### Gradient Ascent

Gradient ascent is an iterative optimization algorithm used to maximize a function by following the direction of steepest ascent of the gradient. The gradient  $\nabla f(x, y)$  at a point  $(x, y)$  points in the direction of the greatest rate of increase of  $f$ . Therefore, moving in the direction of the gradient can lead to local maxima of the function.

### Backtracking Line Search

Backtracking line search is a technique used in optimization to determine an appropriate step size for gradient-based methods. Instead of using a fixed step size, backtracking line search dynamically adjusts the step size based on certain criteria, ensuring that the function value increases sufficiently in the direction of the gradient.

### Armijo Condition

The Armijo condition is a key component of backtracking line search. It ensures that the function value decreases sufficiently with respect to the gradient direction and the chosen step size. The Armijo condition is typically expressed as:

$$f(x + h \cdot g) \leq f(x) + \sigma \cdot h \cdot (g^T g)$$

where  $h$  is the step size,  $g$  is the gradient vector,  $\sigma$  is a small constant (e.g.,  $0 < \sigma < 1$ ), and  $g^T g$  denotes the squared norm of the gradient.

### Armijo Condition vs. Exact Line Search

The Armijo condition is preferred over exact line search methods for several reasons:

- **Computational Efficiency:** Exact line search methods require costly evaluations of the objective function, especially in high-dimensional spaces. In contrast, the Armijo condition only requires a few function evaluations to determine a suitable step size.
- **Robustness:** The Armijo condition provides a balance between ensuring sufficient decrease in the function value and avoiding excessive computation associated with exact line search methods.
- **Adaptability:** Backtracking line search with the Armijo condition can handle noisy or ill-conditioned objective functions more effectively compared to exact line search methods.

## Gradient Ascent Algorithm in MATLAB

Homework #5 required the implementation of a the gradient ascent algorithm with inexact (backtracking) line search to solve a 2D unconstrained optimization problem. Here's a brief explanation of the key components and flow of the code, which can be found in Listing 3:

### Inputs

- `f` – Objective function (anonymous function of `x` and `y`)
- `xi` & `yi` – Initial guesses for `x` and `y`
- `tol` – Error tolerance for convergence
- `sigma` – Armijo condition constant
- `beta` – Backtracking constant

### Outputs

- `xypos` – Array containing the (`x`, `y`) coordinates at each step
- `numsteps` – Number of steps required for convergence
- `numfneval` – Number of function evaluations

### Optimization Process

#### 1. Initialization

- Set initial parameters (`numsteps`, `numfneval`, `maxiter`, `converged`).
- Initialize the current position `x` with the initial guesses.

#### 2. Gradient Ascent with Backtracking Line Search

- Iterate until convergence (`converged = true`) or maximum iterations (`maxiter`) are reached.
- Compute the gradient `g` of the function at the current position `x`.
- Check the convergence condition based on the gradient norm (`err < tol`).
- Implement backtracking line search using the Armijo condition to ensure a sufficient decrease in the function value.
- Update the position `x` based on the computed step size (`h`) and the gradient direction.
- Record the updated position in `xypos`.

#### 3. Contour Plot

- After optimization, visualize the path taken (`xypos`) overlaid on the contour plot of the function.
- Evaluate the function `f` over a grid of `x` and `y` values to create the contour plot.
- Plot the starting point (`xi`, `yi`), the path (`xypos`), and the optimal solution (`xypos(end,1)`, `xypos(end,2)`)

### Utility Functions

- `grad`: Estimates the gradient of `f` using central finite-difference approximation.
- `euclideanNorm`: Computes the Euclidean norm of a vector.
- `applyFigureProperties`: Configures figure properties for plotting.

## Recommendations

- Ensure the objective  $f$  is well-defined and behaves appropriately for the optimization task.
- Tune the parameters ( $\text{tol}$ ,  $\text{sigma}$ ,  $\text{beta}$ ,  $\text{maxiter}$ ) based on the characteristics of the objective function.
- Verify the gradient approximation ( $\text{grad}$ ) accuracy with a smaller perturbation ( $\text{delta}$ ) if needed.
- To find a local minima (i.e. gradient descent), simply negate the sign of the gradient when updating the position at each step, and also modify the Armijo condition to ensure that the function **decreases** by a minimum amount at each step.

## Simulations

Consider the set of input parameters specified for the optimization algorithm using gradient ascent with backtracking line search. Each scenario is defined by different values of tolerance ( $\text{tol}$ ), Armijo condition constant ( $\sigma$ ), and backtracking constant ( $\beta$ ). The objective function  $f(x, y)$  and initial guesses for  $x$  and  $y$  are also specified.

Here is an example of how the inputs are defined for the base case:

```

1 xi = -10; yi = 5;           % initial guesses
2 f = @(x,y) -10*(x-2).^2 - 5*(y+3).^2 + 20; % objective function
3 tol = 0.01;                 % tolerance
4 sigma = 0.0001;             % armijo condition constant
5 beta = 0.5;                 % backtracking constant

```

Listing 1: Function Inputs

## Objective function and Optimization Setup

The objective function  $f(x, y) = -10(x - 2)^2 - 5(y + 3)^2 + 20$  represents a quadratic function with a known maximum. The goal is to find the local maximum of  $f(x, y)$  starting from the initial guesses  $X_i = -10$  and  $y_i = 5$ .

## Scenarios

Table 1 presents ten simulation scenarios with varying parameter values ( $\text{tol}$ ,  $\sigma$ ,  $\beta$ ) and corresponding results from the optimization algorithm.

- **Steps:** Number of iterations required to reach convergence (within the specified tolerance).
- **Function Evaluations:** Total number of function evaluations performed during optimization.

Noteworthy scenarios marked with an asterisk (\*) indicate that a plot was generated and the results are discussed further in the Results section.

Table 1: Ten Scenarios Demonstrating Effects of Varying  $\text{tol}$ ,  $\sigma$ , and  $\beta$ 

Scenario	tol	$\sigma$	$\beta$	Steps	Function evaluations
1	0.01	0.0001	0.5	11	93
2*	0.01	0.0001	0.9	328	8834
3	0.01	0.0001	0.1	48	333
4	0.01	0.0001	0.01	87	521
5	0.01	0.01	0.5	11	93
6	0.01	0.1	0.5	10	85
7*	0.01	0.9	0.5	64	716
8	1e-1	0.0001	0.5	9	76
9	1e-4	0.0001	0.5	15	129
10*	1e-6	0.1	0.5	19	164

## Interpretation of Results

The output from the optimization algorithm (Listing 2) for the base case (scenario 1) provides insights into the optimization process:

- **Optimal Point:** The coordinates  $(x, y)$  of the local maximum identified by the algorithm (shown at the final step if convergence is reached).
- **Iterative Process:** Sequence of points  $(x, y)$  visited during the optimization iterations, captured in the `xypos` array.
- **Algorithm Performance:**
  - **Number of Steps:** Indicates the convergence behavior of the algorithm.
  - **Function Evaluations:** Reflects the computational cost in terms of objective function evaluations.

```
>> [xypos,numsteps,numfneval] = calcMaxStudent(f,xi,yi,tol,sigma,beta);
```

Results:

x	y
-10.0000	5.0000
5.0000	-0.0000
1.2500	-1.8750
2.1875	-2.5781
1.7188	-3.1055
2.0703	-3.0396
1.9824	-3.0148
2.0044	-3.0056
1.9989	-3.0021
2.0016	-2.9995
1.9996	-2.9998

Number of steps = 11

Number of function evaluations = 93

Listing 2: Optimization Algorithm Output

## Discussion of Key Scenarios

In this section, we delve into the results and insights gained from the simulations of scenarios 2\*, 7\*, and 10\*.

### Scenario 2\*

Scenario 2\* utilized a small tolerance ( $\text{tol} = 0.01$ ) and a high backtracking constant ( $\beta = 0.9$ ). This combination resulted in a large number of steps (328) and function evaluations (8834) when compared to scenario 1. The optimization algorithm with these parameters displayed a cautious approach, taking smaller steps towards the local maximum due to the significant backtracking. The plot generated for this scenario, seen in Figure 1, shows a detailed path towards convergence, indicating careful adjustment of the step size at each iteration.

### Scenario 7\*

In scenario 7\*, a higher Armijo condition constant ( $\sigma = 0.9$ ) was employed along with standard tolerance ( $\text{tol} = 0.01$ ) and backtracking constant ( $\beta = 0.5$ ). Despite the larger  $\sigma$ , the algorithm required 64 steps and 716 function evaluations to converge. This behavior suggests that while a higher  $\sigma$  allows larger steps, it may also necessitate more cautious adjustments during the line search process. Figure 2 illustrates the optimization path, showcasing the impact of the armijo condition on step size decisions.

### Scenario 10\*

Scenario 10\* explored a scenario with an extremely small tolerance ( $\text{tol} = 1 \times 10^{-6}$ ) and a higher Armijo condition constant ( $\sigma = 0.1$ ). These parameters facilitated convergence in 19 steps with 164 function evaluations. The small tolerance forced the algorithm to meticulously approach the maximum, while the moderate  $\sigma$  allowed more substantial steps compared to other scenarios with similar tolerances. Figure 3 illustrates the optimization journey, emphasizing the balance between precision and computational efficiency.




Figure 1: Optimization Path for Scenario 2




Figure 2: Optimization Path for Scenario 7




Figure 3: Optimization Path for Scenario 10



## MATLAB Code

```

1 % Written by: Francisco Sanudo
2 % Date: 5/1/24
3 % Updated: 5/8/24
4 %
5 % PURPOSE
6 % calcMaxStudent solves a 2D unconstrained optimization problem by finding
7 % the local maximum of a 2D function f(x,y) using the gradient ascent with
8 % inexact (backtracking) line search method. The Armijo condition is used
9 % to ensure the function increases by a minimum amount at each step.
10 %
11 % INPUTS
12 % - f      : Objective function (anonymous func of x and y)
13 % - xi & yi : Initial guesses for x and y
14 % - tol    : Error tolerance for convergence
15 % - sigma  : Armijo condition constant
16 % - beta   : Backtracking constant
17 %
18 % OUTPUTS
19 % - xypos   : Array that contains the (x,y) coordinates at each step
20 % - numsteps : Number of steps required to achieve the termination criteria
21 % - numfeval : Number of function evaluations
22 % --> Everytime f(x,y) is invoked, a counter variable increases by 1
23 %
24 % EXAMPLE
25 % xi = -10; yi = 5;                % initial guesses
26 % f = @(x,y) -10*(x-2).^2 - 5*(y+3).^2 + 20; % objective function
27 % tol = 1e-2;                      % tolerance
28 % sigma = 0.0001;                  % armijo condition constant
29 % beta = 0.5;                      % backtracking constant
30 %
31 % [xypos,numsteps,numfeval] = calcMaxStudent(f,xi,yi,tol,sigma,beta);
32 %
33 % OTHER
34 % .m files required                : MAIN.m (calling script)
35 % Files required (not .m)          : none
36 % Built-in MATLAB functions used   : numel, zeros
37 % Utility functions                 : grad, eucildeanNorm, applyFigureProperties, myMax
38 %
39 % REFERENCES
40 % Optimization (notes), P. Nissenon
41 %
42 %
43 function [xypos,numsteps,numfeval] = calcMaxStudent(f,xi,yi,tol,sigma,beta)
44
45 % Initialize Variables
46 numsteps = 1;                % Set the iteration counter
47 numfeval = 0;                % Set the function evaluations counter
48 maxiter = 10000;             % Maximum iterations
49 converged = false;           % Convergence condition
50
51 X = [xi; yi];                % Set initial condition

```

```
52 xypos = X'; % Initialize position history
53
54 %% Gradient ascent with inexact (backtracking) line search method
55
56 % Iterate
57 while ~converged && numsteps < maxiter
58     % Compute gradient
59     g = grad(f, X);
60     numfneval = numfneval + 4; % Increment function evaluation counter
61
62     % Compute error using gradient norm
63     err = euclideanNorm(g);
64
65     if err < tol
66         converged = true;
67     else
68         % Backtracking line search
69         h = 1; % Initial step size
70
71         % Temporary step
72         xtemp = X(1) + g(1)*h;
73         ytemp = X(2) + g(2)*h;
74
75         % Function value at current step
76         fcurrent = f(X(1),X(2));
77         numfneval = numfneval + 1; % Increment function evaluation counter
78
79         % Armijo Condition
80         while f(xtemp, ytemp) - fcurrent < sigma*(g'*g)*h
81             % Reduce step size using backtracking
82             h = beta * h;
83
84             % Increment function evaluation counter
85             numfneval = numfneval + 1;
86
87             % Update temporary variables
88             xtemp = X(1) + g(1)*h;
89             ytemp = X(2) + g(2)*h;
90         end
91
92         % Update position
93         X = X + h*g;
94
95         % Update position history
96         xypos = [xypos; X'];
97
98         % Increment iteration counter
99         numsteps = numsteps + 1;
100     end
101 end
102
103 % Increment function evaluation counter (to account for final step when
104 % convergence is reached)
```

```
105 numfneval = numfneval + 1;
106
107 %% Contour plot that shows the path taken to the maximum
108
109 % Extract x and y coordinates from xypos
110 x_coords = xypos(:, 1);
111 y_coords = xypos(:, 2);
112
113 % Calculate maximum magnitudes of x and y coordinates
114 max_x_magnitude = myMax(abs(x_coords));
115 max_y_magnitude = myMax(abs(y_coords));
116
117 % Calculate x and y ranges (20% larger than max magnitude)
118 x_range = [-1, 1] * (1.2 * max_x_magnitude);
119 y_range = [-1, 1] * (1.2 * max_y_magnitude);
120
121 % Calculate grid spacing based on the modified ranges
122 dx = abs((x_range(end) - x_range(1))) / 100;
123 dy = abs((y_range(end) - y_range(1))) / 100;
124
125 % Define grid for contour plot
126 x = x_range(1):dx:x_range(2);
127 y = y_range(1):dy:y_range(2);
128
129 % Evaluate the function at each point on grid
130 % Each row corresponds to a constant y value and each column corresponds
131 % to a constant x value. This means that z(i, j) represents the function
132 % value at x(j) and y(i)
133
134 z = zeros(numel(y),numel(x));
135
136 for i = 1:numel(x)
137     for j = 1:numel(y)
138         z(i,j) = f(x(j), y(i));
139     end
140 end
141
142 % Create figure and apply figure properties
143 f = figure;
144 position = [0.2, 0.2, 0.5, 0.6];
145 applyFigureProperties(f, position)
146
147 hold on; % plot over same axes
148 ax = gca; % get current axes
149
150 % Plot contours
151 contour(x,y,z,'ShowText','on');
152
153 % Plot the path to the optimal solution
154 plot(x_coords,y_coords,'k')
155 plot(x_coords,y_coords,'ko','MarkerSize',3,'MarkerFaceColor','k')
156
157 % Plot starting point
```

```

158 plot(xi,yi,'bo','Markersize',6,'MarkerFaceColor','b')
159
160 % Plot the optimal solution
161 plot(x_coords(end),y_coords(end),'mo','Markersize',6,'MarkerFaceColor','m')
162
163 % Axis properties
164 set(ax,'TickLabelInterpreter','latex')
165 title('Gradient Ascent w/ Armijo Condition')
166 xlabel('$x$')
167 ylabel('$y$')
168 legend('', 'Path', '', '$x_i$, $y_i$', '$x^*$, $y^*$')
169 grid on
170 % axis equal
171
172 % Colormap & colorbar properties
173 colormap(ax,'turbo');
174 cb = colorbar;
175 cb.TickLabelInterpreter = 'latex';
176 title(cb,'$\rm Level$', 'Interpreter','latex')
177
178 end
179
180 %-----
181
182 function g = grad(f,X)
183 % grad estimates the gradient of the function f using a centered finite-
184 % difference approximation.
185 %
186 % Input:
187 %   f: Anonymous function representing the objective function f(x, y)
188 %   X: Input vector containing the current point (x, y)
189 %
190 % Output:
191 %   g: Estimated gradient vector of f at the point X
192
193 n = numel(X);           % Dimension of input vector X
194 g = zeros(n, 1);        % Initialize gradient vector
195 delta = 0.00001;        % Perturbation parameter
196
197 % Estimate gradient
198 for i = 1:n
199     % Create a copy of X to perturb
200     X_perturbed = X;
201
202     % Perturb the i-th element
203     X_perturbed(i) = X_perturbed(i) + delta;
204
205     % Calculate the forward difference
206     f_forward = f(X_perturbed(1),X_perturbed(2));
207
208     % Perturb in the negative direction
209     X_perturbed(i) = X_perturbed(i) - 2 * delta;
210

```

```
211     % Calculate the backward difference
212     f_backward = f(X_perturbed(1),X_perturbed(2));
213
214     % Estimate partial derivative w.r.t. i-th element using central
215     % difference formula
216     g(i) = (f_forward - f_backward) / (2 * delta);
217 end
218
219 end
220
221 %-----
222
223 function euclidean_norm_value = euclideanNorm(vec)
224 % euclideanNorm computes the Euclidean norm (magnitude) of an n-dimensional vector.
225 %
226 % Input:
227 %   vec: An n-dimensional vector (column vector)
228 %
229 % Output:
230 %   euclidean_norm_value: The Euclidean norm of the input vector 'vec'
231
232 % Square each element of vector
233 squared_elements = vec.^2;
234
235 % Initialize sum_of_squares
236 sum_of_squares = 0;
237
238 % Loop through each element
239 for i = 1:numel(vec)
240     % Square the current element
241     squared_element = vec(i)^2;
242
243     % Add the squared element to the sum_of_squares
244     sum_of_squares = sum_of_squares + squared_element;
245 end
246
247 % Compute the Euclidean norm
248 euclidean_norm_value = sqrt(sum_of_squares);
249
250 end
251
252 %-----
253
254 function [maximum, idx] = myMax(A)
255 % myMax calculates the largest element of an array A and returns its value
256 % along with the index of that element in the array.
257 %
258 % Input:
259 %   A: Array of numeric values
260 %
261 % Output:
262 %   maximum: Largest element in the array A
263 %   idx: Index of the largest element in the array A
```

```
264
265 % Initialize variables to track the maximum value and its index
266 maximum = A(1); % Assume the first element is the maximum initially
267 idx = 1;       % Index of the assumed maximum
268
269 % Iterate through the array to find the actual maximum value and its index
270 for j = 2:numel(A)
271     if A(j) > maximum
272         % Update the maximum value and its index if a larger value is found
273         maximum = A(j);
274         idx = j;
275     end
276 end
277
278 end
279
280 %-----
281
282 function applyFigureProperties(figHandle, position)
283 % applyFigureProperties sets specific properties for a given figure handle.
284 %
285 % Inputs:
286 %   - 'figHandle': Handle to the figure (obtained using 'figure' or 'gcf')
287 %   - 'position': Position vector [left, bottom, width, height] in normalized units
288 %                 specifying the figure's position and size
289 % Output:
290 %   None (modifies the specified figure properties directly)
291 %
292 % Example Usage:
293 %   applyFigureProperties(fig, [0.2, 0.2, 0.5, 0.6]);
294
295 set(figHandle, ...
296     'Units', 'normalized', ...
297     'Position', position, ...
298     'DefaultTextInterpreter', 'latex', ...
299     'DefaultLegendInterpreter', 'latex', ...
300     'DefaultAxesFontSize', 14);
301
302 end
```

Listing 3: Gradient Ascent Algorithm with Inexact (Backtracking) Line Search for 2D Optimization